

# Unit tests for Linux kernel hackers

How to avoid hours of painful debugging

Mihir Mehta

Systems Core Group  
Samsung Research Institute - Noida

`mihir.mehta@samsung.com`

February 3, 2014

# Overview

Overview of the material.

- ▶ What are unit tests?
- ▶ Why unit tests?
- ▶ Accessing unit test results
- ▶ Summary

# What are unit tests?

# What are unit tests?

- ▶ Ideally, a unit test involves calling a *single* function and *verifying* its return value against an *expected* value.
- ▶ Running a test can have two outcomes.
  1. Pass: Congrats! Your function works correctly on this test input!
  2. Fail: This again raises multiple outcomes:
    - 2.1 Function body is wrong. (bad, but manageable if you know what you're doing.)
    - 2.2 Expected value is wrong (indicates you don't know what you're doing, i. e. very bad.)

# Why unit tests?

- ▶ Central component of test-driven design.
- ▶ Provide convenient way to verify our code.
- ▶ Save manual labour (i.e. less time at desk!)
- ▶ Crucial: help prevent regressions.

# Accessing unit test results

- ▶ In a userspace program, (i.e. what we're used to) it's simple to make a separate test executable and link the functions we are testing into this executable.
- ▶ This way, we maintain a logical separation of the test code from the codebase itself.
- ▶ We also avoid bloating our executables.
- ▶ Making a separate test executable is not possible within kernel code - why?

# Accessing unit test results

- ▶ Making a separate test executable is not possible within kernel code - why?
  - ▶ Can't be a userspace program - userspace programs can't link to kernel code (privilege levels and other issues)
  - ▶ Can't be a kernel - how are you going to run two kernels at once?
- ▶ So, how do we implement unit tests for kernel code now?

# Accessing unit test results

- ▶ So, how do we implement unit tests for kernel code now?
  - ▶ Kernel modules!
- ▶ With modules, we can easily access kernel functions and test them.
- ▶ We can run all the tests at once, in the init function of the module, after setting a load priority for the module.
- ▶ If we use loadable modules (instead of static modules), we can avoid kernel bloat too - simply unload the module when done.
- ▶ One detail remains - how do we access the results of these tests in userspace?



## Accessing unit test results

- ▶ How do we access the results of our unit tests in userspace?
  - ▶ Special kernel filesystems - sysfs and debugfs
- ▶ It's actually quite common for kernel variables to need to be readable and possibly modifiable from userspace. For this purpose, the Linux kernel provides certain special file systems.
  - ▶ procfs - NOT to be used for debugging.
  - ▶ sysfs - should not be used for debugging (but we do it anyway).
  - ▶ debugfs - explicitly intended for debugging.
- ▶ sysfs - used by many modules which need to communicate with user space; also abused for debugging; has strict restrictions on the types of exported values (only one numeric value per file).
- ▶ debugfs - used by many modules for debugging; has no restrictions on the types of exported values.
- ▶ Both of the above are sufficiently widely used to appear in nearly all distribution kernels (including ours).
- ▶ sysfs is generally mounted at /sys, while debugfs is generally mounted at /sys/kernel/debug.

# Using debugfs for debugging

- ▶ Every kernel module should store its debugging files in a separate directory in debugfs. Optionally, subdirectories can be used for better organisation.

- ▶ This function is used for creating a directory -

```
struct dentry *debugfs_create_dir(const char  
*name, struct dentry *parent);
```

If the value of `parent` is `NULL`, then the directory is created at the debugfs root.

- ▶ The `struct dentry *` pointer obtained from this function can now be used to create files in the directory, using this function -

```
struct dentry *debugfs_create_file(const char  
*name, umode_t mode, struct dentry *parent, void  
*data, const struct file_operations *fops);
```

## Using debugfs for debugging

- ▶ Usually, though, the function `debugfs_create_file` is too bulky for our purposes - we'd rather not have to make a set of file operations.
- ▶ The `debugfs` module also provides a few convenience functions for common debugging use cases, to remove the need for making sets of file operations.
- ▶ In cases where we only want to export a numeric value, we can take advantage of these functions:
  - ▶ `struct dentry *debugfs_create_u8(const char *name, umode_t mode, struct dentry *parent, u8 *value);`
  - ▶ `struct dentry *debugfs_create_u16(const char *name, umode_t mode, struct dentry *parent, u16 *value);`
  - ▶ `struct dentry *debugfs_create_u32(const char *name, umode_t mode, struct dentry *parent, u32 *value);`
  - ▶ `struct dentry *debugfs_create_u64(const char *name, umode_t mode, struct dentry *parent, u64 *value);`

# Acknowledgements

- ▶ `doc/Documentation/filesystems/debugfs.txt` in the kernel source. (primary reference, authored by Jonathan Corbet.)