

Unit tests for Linux kernel hackers

How to avoid hours of painful debugging

Mihir Mehta

Systems Core Group
Samsung Research Institute - Noida

`mihir.mehta@samsung.com`

February 3, 2014

Overview

Overview of the material.

- ▶ What are unit tests?
- ▶ Why unit tests?
- ▶ Accessing unit test results
- ▶ Example
- ▶ Summary

What are unit tests?

What are unit tests?

- ▶ Ideally, a unit test involves calling a *single* function and *verifying* its return value against an *expected* value.
- ▶ Running a test can have two outcomes.
 1. Pass: Congrats! Your function works correctly on this test input!
 2. Fail: This again raises multiple outcomes:
 - 2.1 Function body is wrong. (bad, but manageable if you know what you're doing.)
 - 2.2 Expected value is wrong (indicates you don't know what you're doing, i. e. very bad.)

Why unit tests?

- ▶ Central component of test-driven design.
- ▶ Provide convenient way to verify our code.
- ▶ Save manual labour (i.e. less time at desk!)
- ▶ Crucial: help prevent regressions.

Accessing unit test results

- ▶ In a userspace program, (i.e. what we're used to) it's simple to make a separate test executable and link the functions we are testing into this executable.
- ▶ This way, we maintain a logical separation of the test code from the codebase itself.
- ▶ We also avoid bloating our executables.
- ▶ Making a separate test executable is not possible within kernel code - why?

Accessing unit test results

- ▶ Making a separate test executable is not possible within kernel code - why?
 - ▶ Can't be a userspace program - userspace programs can't link to kernel code (privilege levels and other issues)
 - ▶ Can't be a kernel - how are you going to run two kernels at once?
- ▶ So, how do we implement unit tests for kernel code now?

Accessing unit test results

- ▶ So, how do we implement unit tests for kernel code now?
 - ▶ Kernel modules!
- ▶ With modules, we can easily access kernel functions and test them.
- ▶ We can run all the tests at once, in the init function of the module, after setting a load priority for the module.
- ▶ If we use loadable modules (instead of static modules), we can avoid kernel bloat too - simply unload the module when done.
- ▶ One detail remains - how do we access the results of these tests in userspace?

Accessing unit test results

- ▶ How do we access the results of our unit tests in userspace?
 - ▶ Special kernel filesystems - sysfs and debugfs
- ▶ It's actually quite common for kernel variables to need to be readable and possibly modifiable from userspace. For this purpose, the Linux kernel provides certain special file systems.
 - ▶ procfs - NOT to be used for debugging.
 - ▶ sysfs - should not be used for debugging (but we do it anyway).
 - ▶ debugfs - explicitly intended for debugging.
- ▶ sysfs - used by many modules which need to communicate with user space; also abused for debugging; has strict restrictions on the types of exported values (only one numeric value per file).
- ▶ debugfs - used by many modules for debugging; has no restrictions on the types of exported values.
- ▶ Both of the above are sufficiently widely used to appear in nearly all distribution kernels (including ours).
- ▶ sysfs is generally mounted at /sys, while debugfs is generally mounted at /sys/kernel/debug.

Using debugfs for debugging

- ▶ Every kernel module should store its debugging files in a separate directory in debugfs. Optionally, subdirectories can be used for better organisation.
- ▶ This function is used for creating a directory -

```
struct dentry *debugfs_create_dir(const char *name, struct dentry *parent);
```

If the value of `parent` is `NULL`, then the directory is created at the debugfs root.

- ▶ The `struct dentry *` pointer obtained from this function can now be used to create files in the directory, using this function -

```
struct dentry *debugfs_create_file(const char *name, umode_t mode, struct dentry *parent, void *data, const struct file_operations *fops);
```

Using debugfs for debugging

- ▶ Usually, though, the function `debugfs_create_file` is too bulky for our purposes - we'd rather not have to make a set of file operations.
- ▶ The `debugfs` module also provides a few convenience functions for common debugging use cases, to remove the need for making sets of file operations.
- ▶ In cases where we only want to export a numeric value, we can take advantage of these functions:
 - ▶ `struct dentry *debugfs_create_u8(const char *name, umode_t mode, struct dentry *parent, u8 *value);`
 - ▶ `struct dentry *debugfs_create_u16(const char *name, umode_t mode, struct dentry *parent, u16 *value);`
 - ▶ `struct dentry *debugfs_create_u32(const char *name, umode_t mode, struct dentry *parent, u32 *value);`
 - ▶ `struct dentry *debugfs_create_u64(const char *name, umode_t mode, struct dentry *parent, u64 *value);`

Example

- ▶ Suppose you make a kernel function
`unsigned int lcs_length(char *s1, char *s2);`
for determining the length of the longest common substring of the two null-terminated strings, `s1` and `s2`.
- ▶ You'll need to test it with a few examples.
 - ▶ `lcs_length("hammer", "ship")` should be 0.
 - ▶ `lcs_length("hammer", "boat")` should be 1.
 - ▶ `lcs_length("sledgehammer", "jackhammer")` should be 6.
- ▶ The more complex the function, the more test cases it requires.
- ▶ How do we do this?

Example

- ▶ All our code will go into our module init function. Since we are building a static module, there is no need for a module exit function.

- ▶ Start by making a directory at the root of debugfs:

```
struct dentry *dir =  
debugfs_create_dir("lcs_test", NULL);
```

- ▶ Next, we create three files inside this directory:

```
struct dentry *test01 =  
debugfs_create_u32("test01", 0_RDONLY, dir,  
(u32)(lcs_length("hammer", "ship") == 0));  
struct dentry *test02 =  
debugfs_create_u32("test02", 0_RDONLY, dir,  
(u32)(lcs_length("hammer", "boat") == 1));  
struct dentry *test03 =  
debugfs_create_u32("test03", 0_RDONLY, dir,  
(u32)(lcs_length("sledgehammer", "jackhammer") ==  
6));
```

Example

- ▶ After recompiling the kernel with our new static module, we can check our results.

```
cat $(SYSFS_PATH)/lcs_test/test*
```

This will display 3 numbers. If none of them are zero, we are done.

Acknowledgements

- ▶ `doc/Documentation/filesystems/debugfs.txt` in the kernel source. (primary reference, authored by Jonathan Corbet.)